

# HydraICN

## Scalable Content Exchange in Challenged ICNs

Bilal Gill


Jan Seedorf

Dirk Kutscher

NEC Laboratories Europe


Supported by GreenICN Project

# ICN Information Availability in Fragmented Networks



ICN nodes: store-and-forward mode

# Data Mules Between Fragmented Communities



# Challenges

- Achieving „**optimal**“ **data availability** in fragmented ICNs with intermittent, unpredictable connectivity
- **Allocating storage and networking resources** accordingly
- **Leveraging knowledge about data popularity/relevance** to optimize resource usage
- **No global view on popularity**: need good enough decentralized estimation
- **Estimation algorithm**: balance effectiveness and complexity

# Objectives

## **Distributed counting/aggregating of interests**

- Interests serve as popularity counters
- Aggregating interest information without losing too much information

## **Robust protocol**

- Scalable with respect to network size and network lifetimes
- Loop-free to accomodate random, unpredictable movements of ICN nodes

# Naïve Approach: Counting Interests

**Append nonce to each unique Interest message**  
and accumulate nonces in the network

- Interest + nonce == unique interest

**Advantage: can lead to accurate representation**  
of popularity per Interest at many nodes

**Disadvantage: Scalability problem**

- Need to store all nonces per Interest
- Also: need to exchange complete nonce list when two nodes meet

# Our Approach


- Idea: **Aggregate [nonce:count] tuples**
  - Approximate content popularity
- Append nonce to each unique Interest message and **count nonces as a popularity counter: [nonce:count] tuples**
- Data mules maintain sorted list of [nonce:count]** per interest for scalability
- When two data mules meet, they **exchange their Interests**
  - Interests & Popularity estimation gets distributed in network

# Our Approach

- For each Interest, aggregate **[nonce:count]** tuples as follows:
  - Compare** ([nonce1, count1] , [nonce2, count2])
  - IF** nonce1 == nonce2
    - new\_count = MAX(count1, count2) (at both nodes)
  - IF** nonce1 != nonce2
    - New\_nonce = nonce with largest counter([nonce1, count1], [nonce2, count2])
    - New\_count = count1 + count2




# Sample Exchange 1



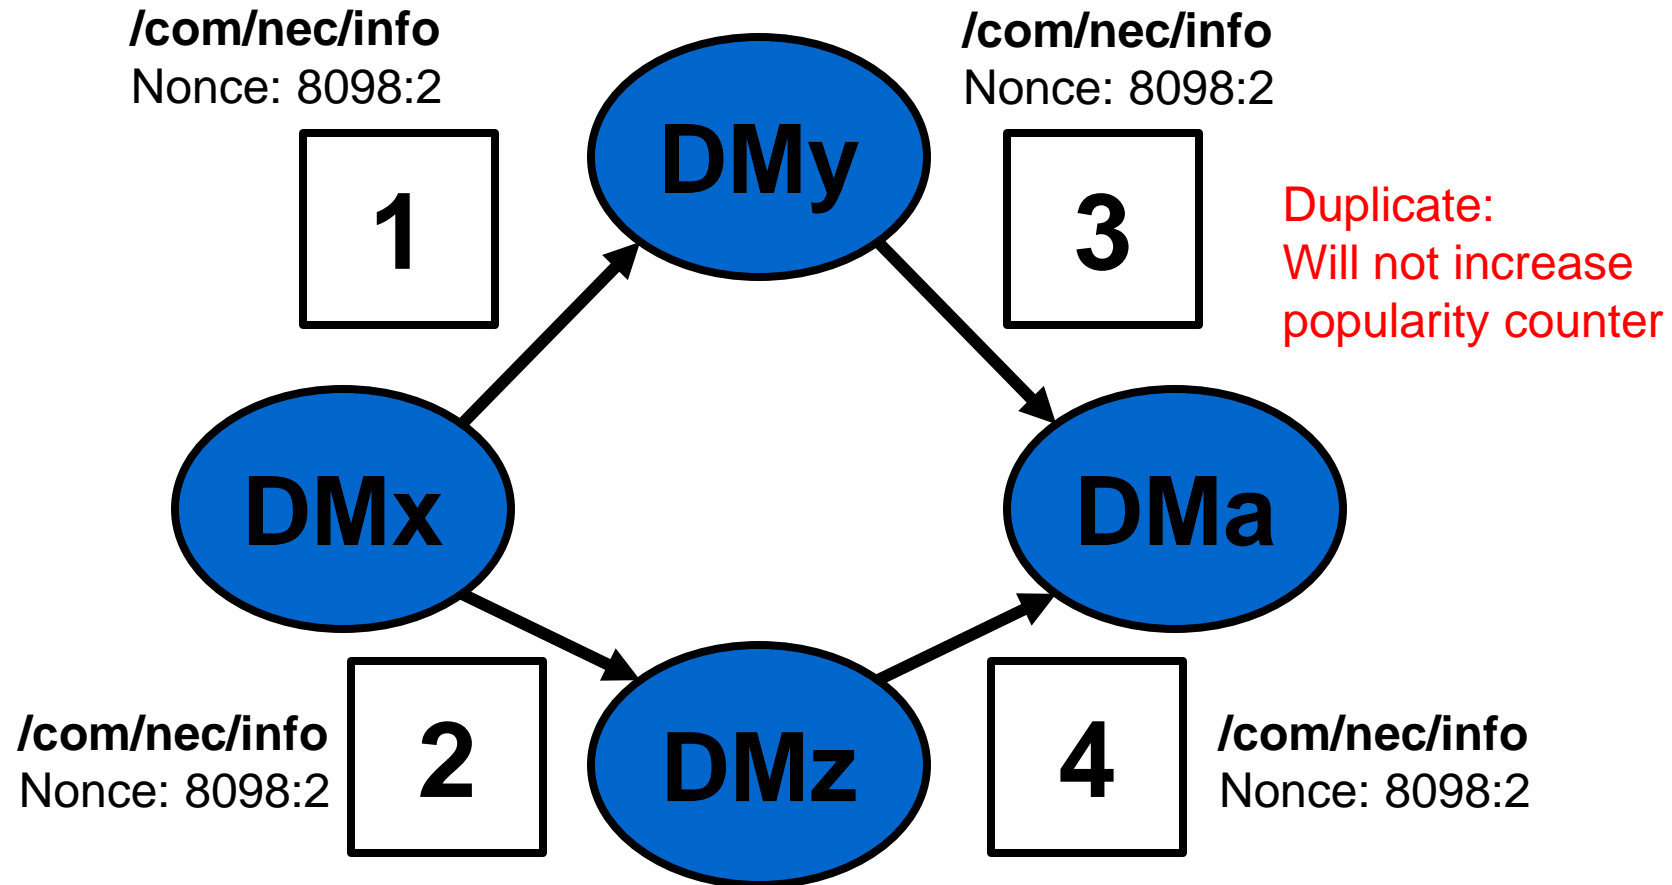
Two end users request the same name from DMx

# Sample Exchange 2



End users are assigned a new nonce with count 2 for the same prefix

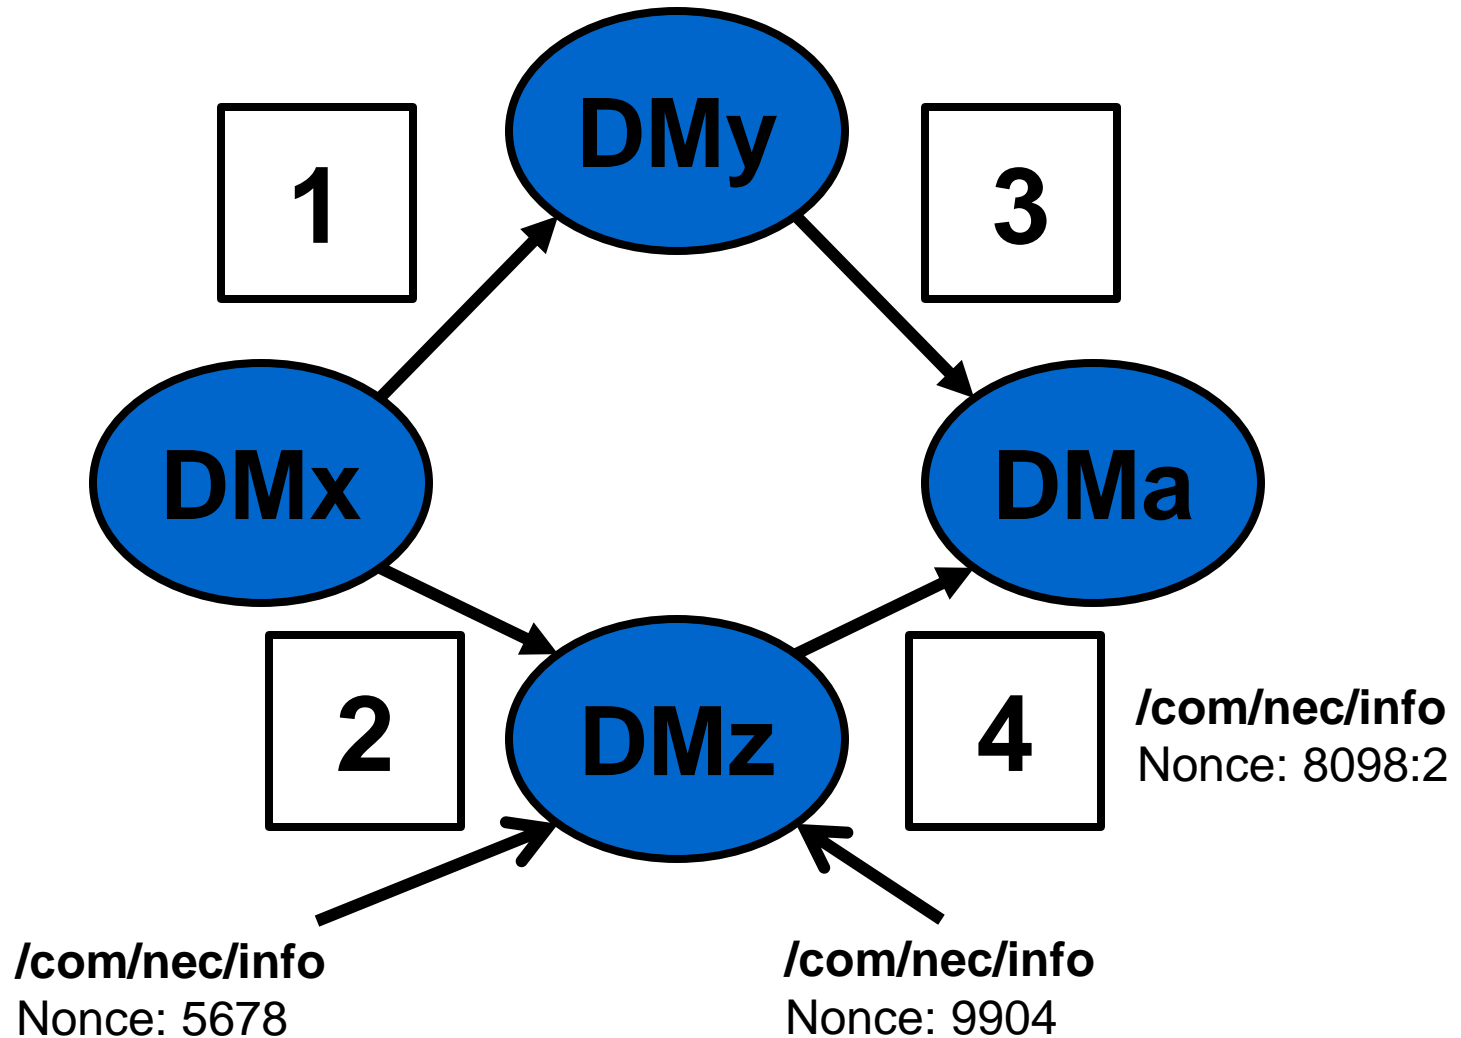
# Sample Exchange 3



DMa receives duplicate query (loop) from DMz,  
thus drops the message (loop detection)

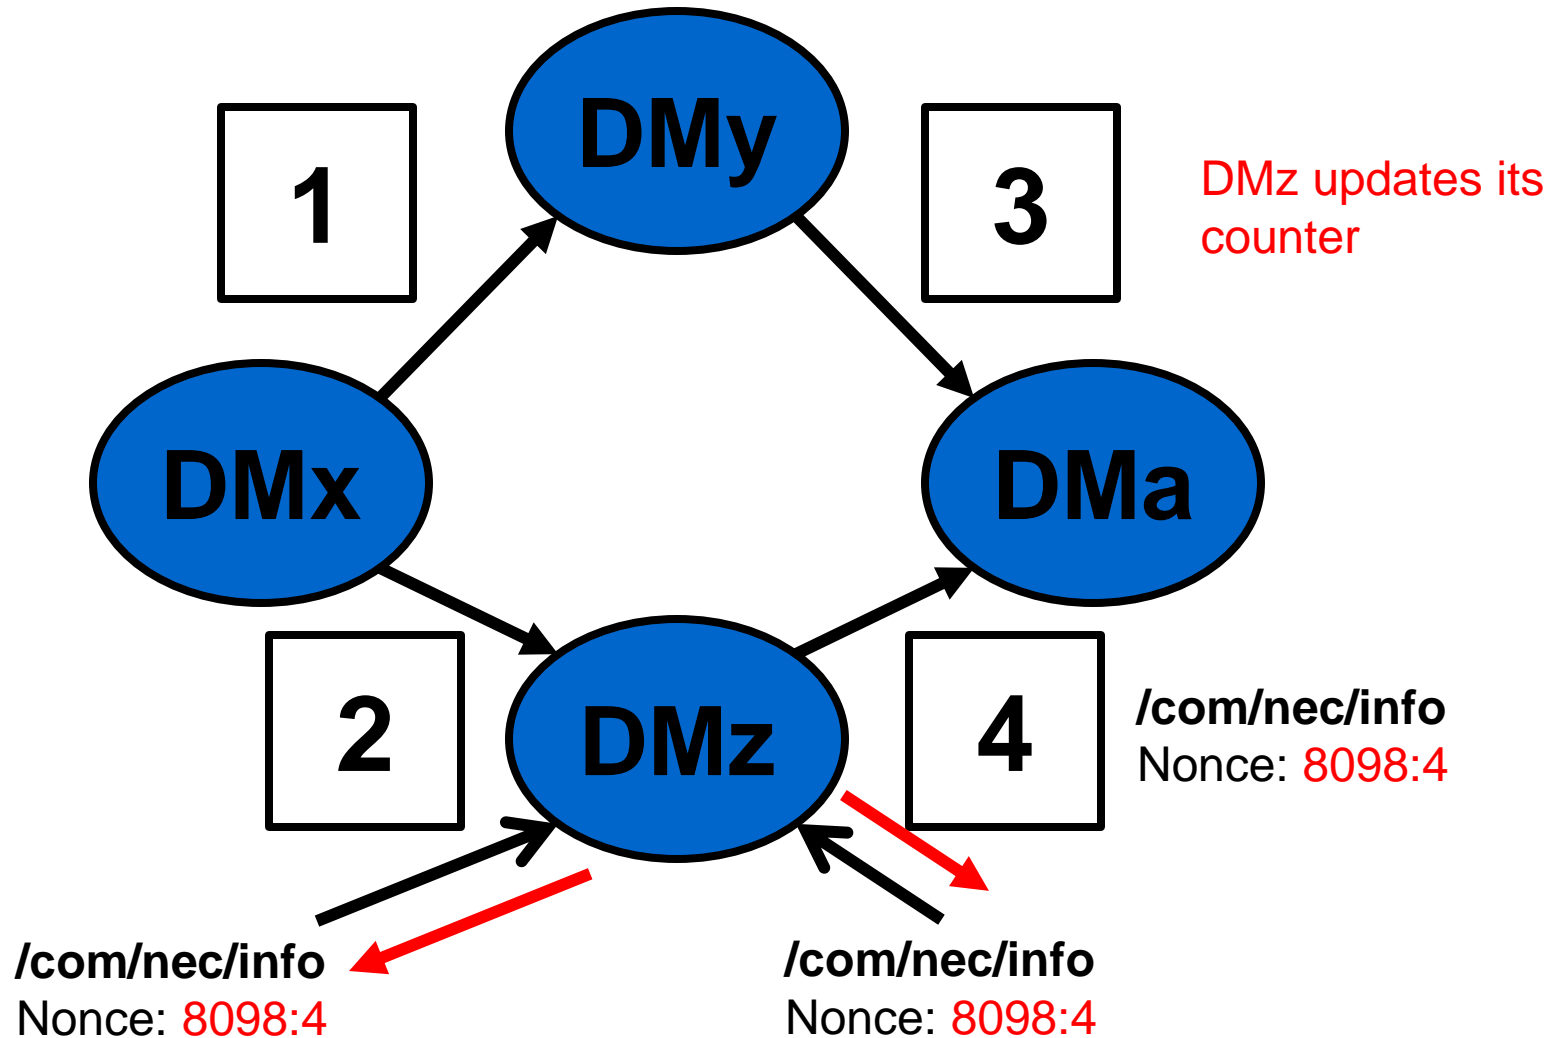
# Sample Exchange 4

Two new user requests with same prefix from DMz



# Sample Exchange 5


DMz updates its count and sends new once to the end user



# Loop Problem



**/com/nec/info**

Nonce: 8098:3



**/com/nec/info**


Nonce: 8098:1




**/com/nec/info**

Nonce: 6789:8

# Loop Problem




# Loop Problem





# Loop Problem

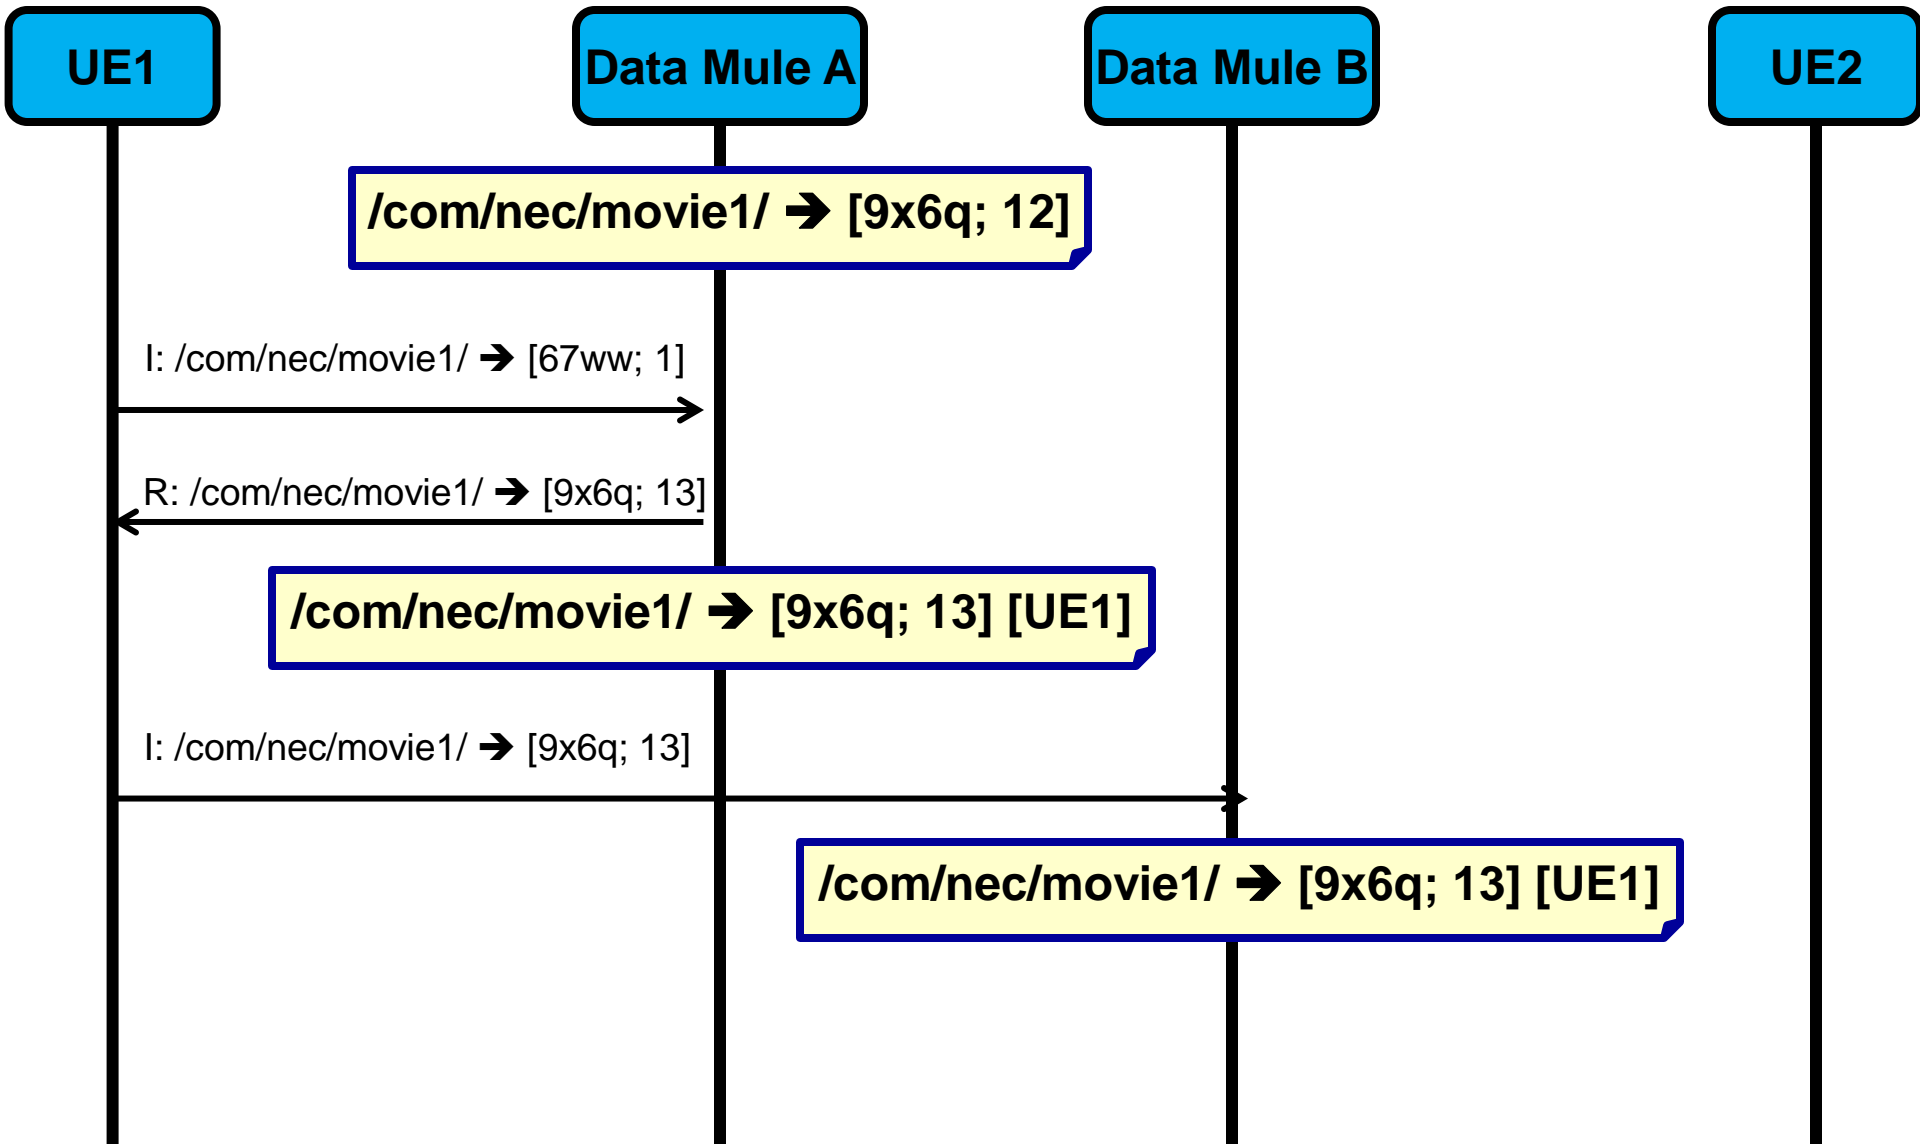


Over-estimating  
Popularity!

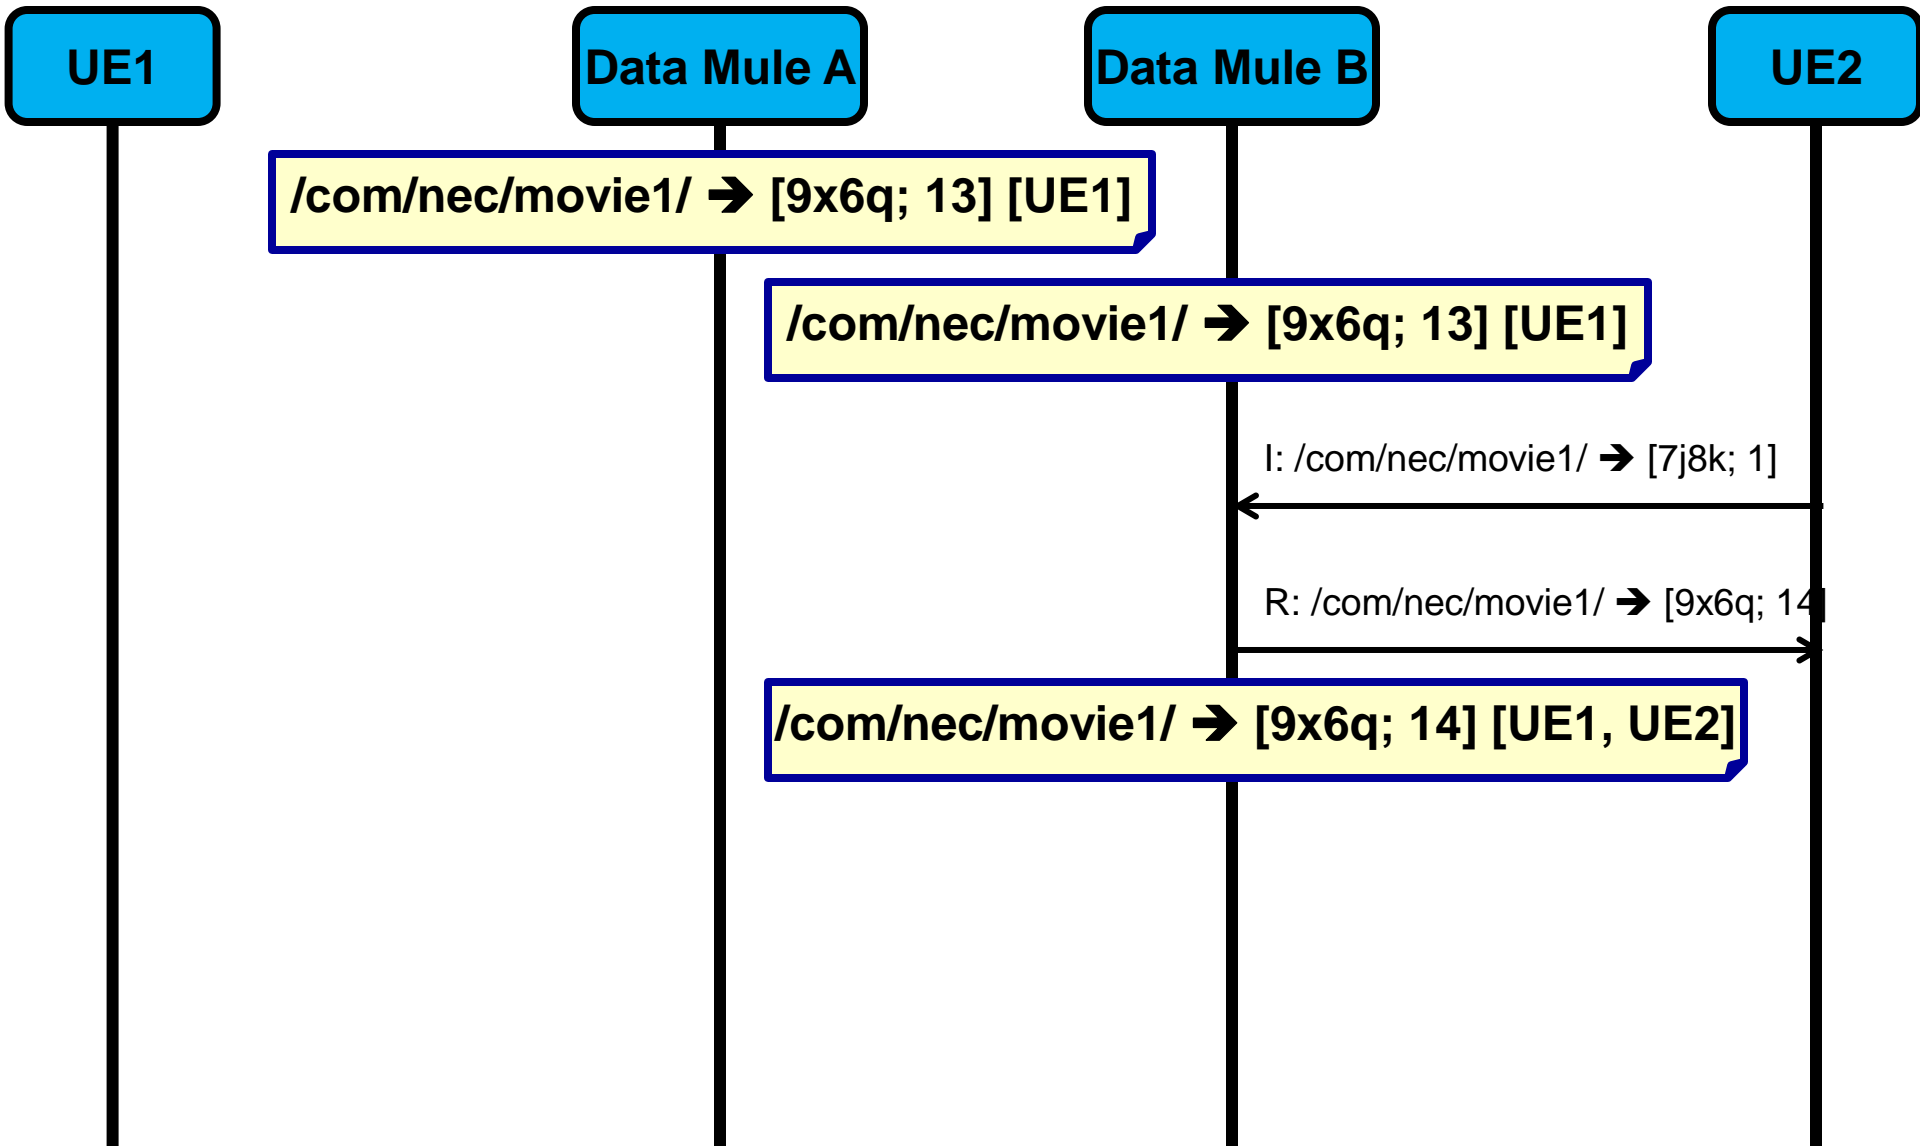
# Loop Prevention

- Data mules keep list of recently encountered mules per Interest
- If a data mule is encountered shortly after the first encounter, counters are not added
- Instead, max-counter rules is applied at both sides
  - `nonce1 != nonce2 (AND recently met)`
    - `New_nonce = nonce with largest counter([nonce1, count1], [nonce2, count2])`
    - `New_count = MAX(count1, count2)`
- Sliding approach: FIFO list of encountered nodes has limited (configurable) size
  - Can trade off accuracy against memory conservation

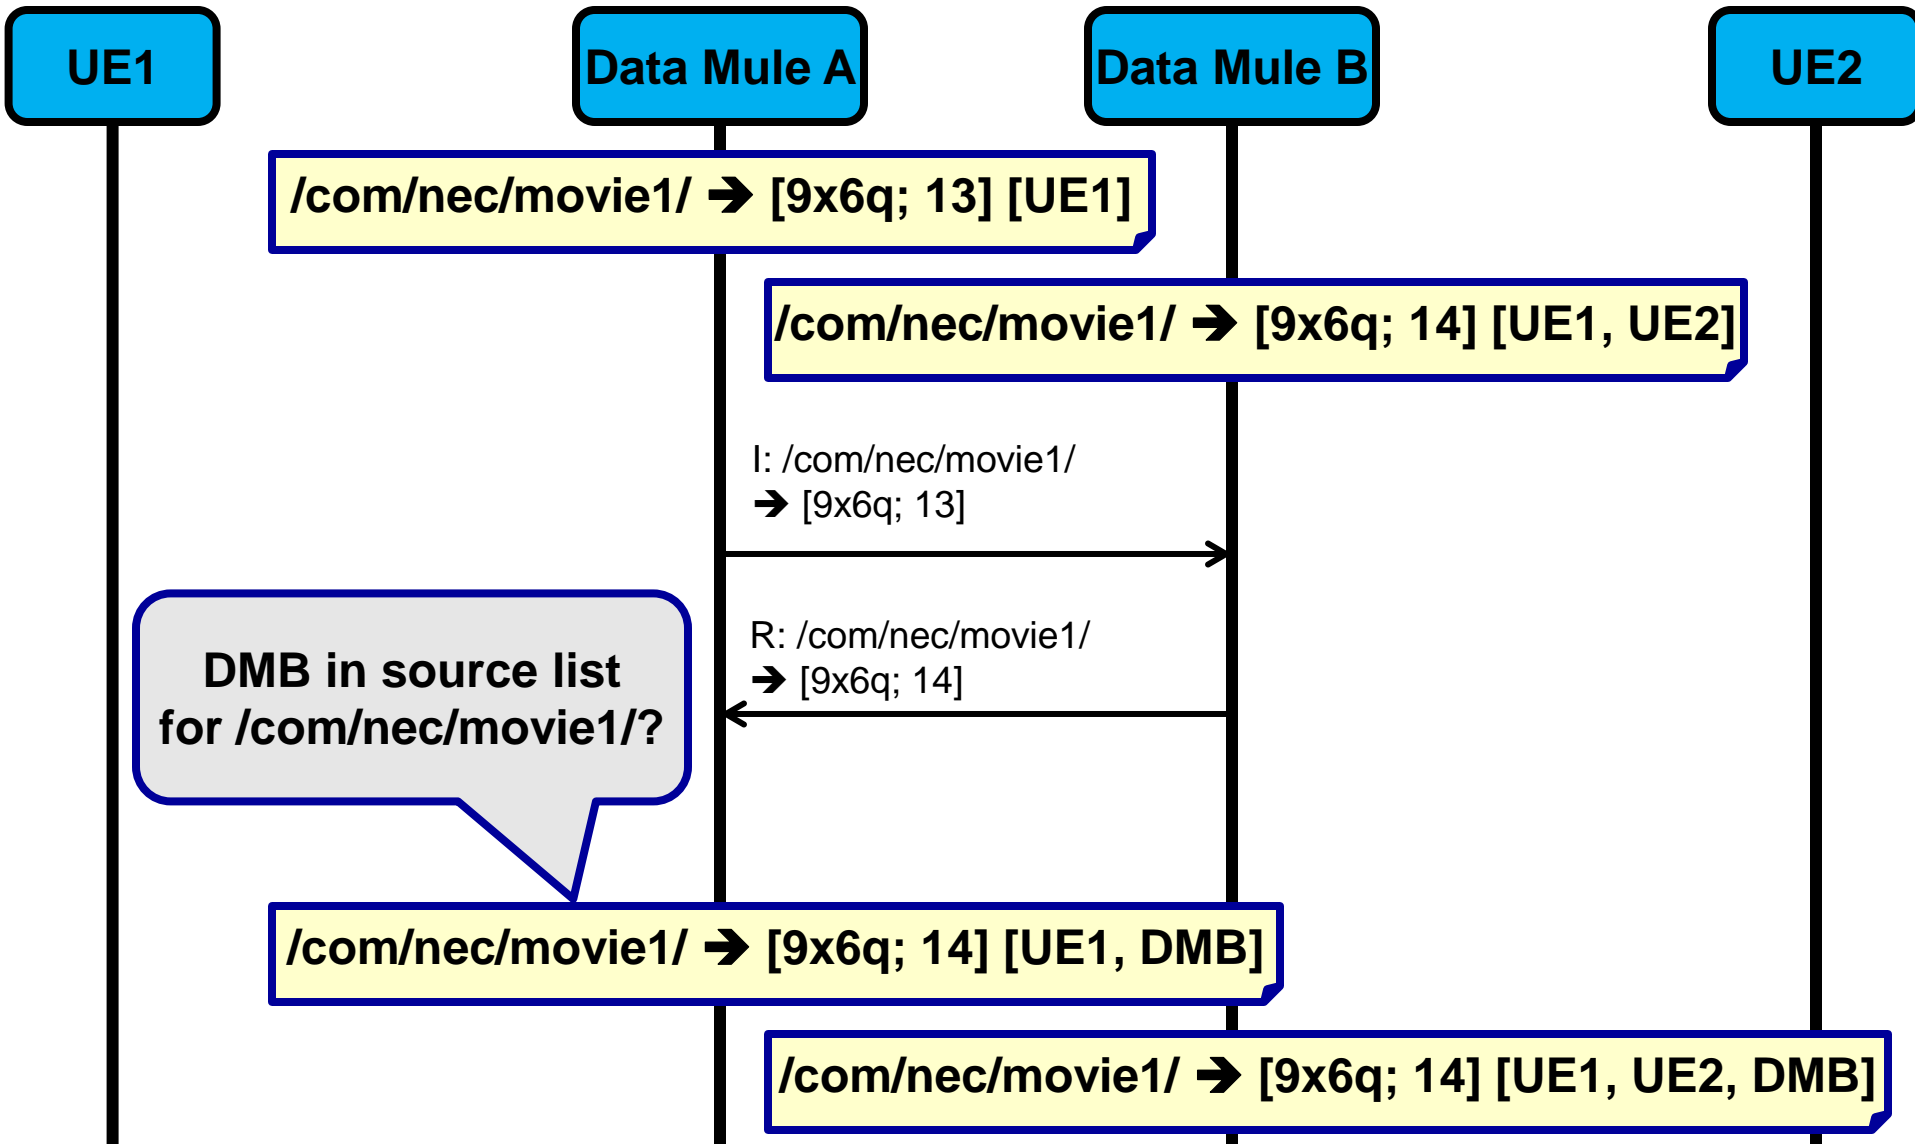
# Exchange Scheme (1)



# Exchange Scheme (2)



# Exchange Scheme (3)



# Implementation

- Interests are retransmitted whenever two nodes meet
- Application on top of CCN creates FIB entries for all the interests (names) we want to send to the next node
- The life time of an interest packet is set to infinite
  - Objective: don't expire interest
  - PITs can become large...
- Current exchange protocol implemented in TCP (not in CCNx protocol)
- Popularity count for all the prefixes in the PIT
  - Colon separates nonce and count

/com/nec/info

8098:4

# Name/Popularity-Based Prioritization

- Contact time between two data mules could be short
  - Data prioritization scheme needed to optimize exchanges
  - Here: **Leverage popularity estimation for prioritization**

## ■ *Name-Based Prioritization protocol (NBP)*

1. Node contact: nodes exchange meta data about cached objects
  - List of names, assessed popularity
2. Requesting node generates Interest packets, ordered by „priority“

- Different categorization schemes possible
  - GreenICN: prioritizing by name / prefix (disaster scenario)


# Evaluation

- Objective: evaluate accuracy of distributed popularity estimation
- Two Fragmented Communities
- 100 distinct data objects
  - Queried from three different DMs in each of the fragmented communities
- Data mules meet each other in a random manner and exchange interests
- Zipf distribution for Interests with  $0.8 < \alpha < 0.9$



# Scenario Diagram

- 3 Nodes per FG which fetch interests from 600 end users
- Different sets of names (popular content) in FGs



Fragmented Community 1

Fragmented Community 2

# Measurement Setup

## Areas with data mules

- Emulation: Mules meet randomly, no disruption during intermeeting times

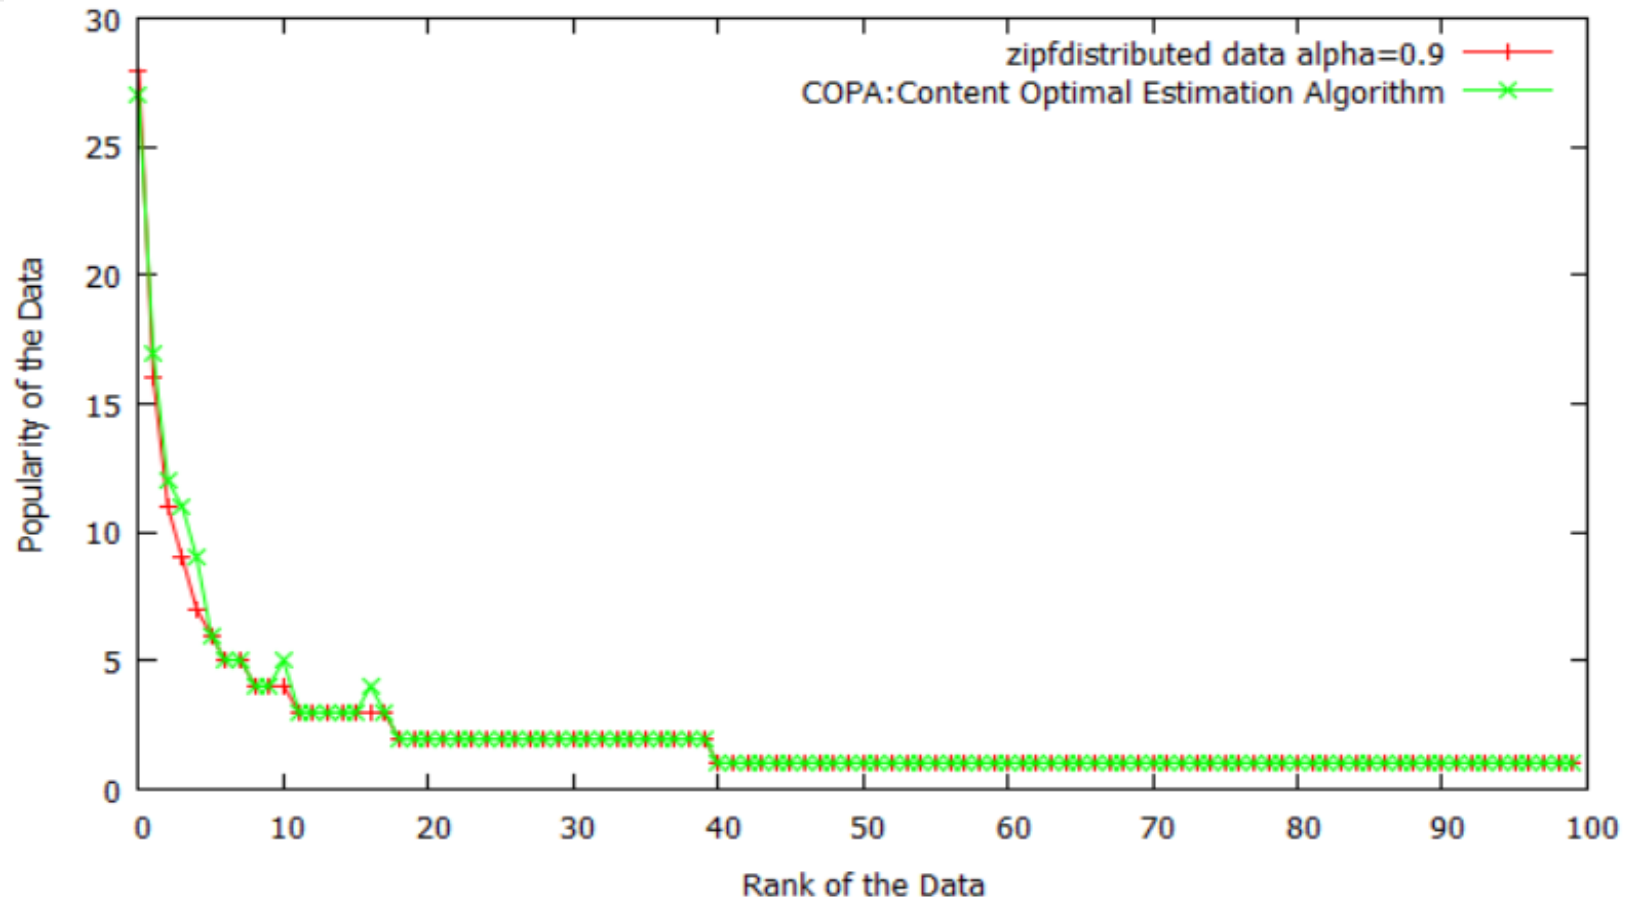
## Details of the test setup

- 6 nodes, 2 Fragmented Communities with different set of names
- Three nodes in their respective FG fetch interest from end users
- Meet DM at random times
- 15 contacts between different DM at different times
- Object size: 8 KB
- Approx. 500 Interests issued by users

## Test specs

- 10 runs for each of the normalized results
- Each runs takes several minutes

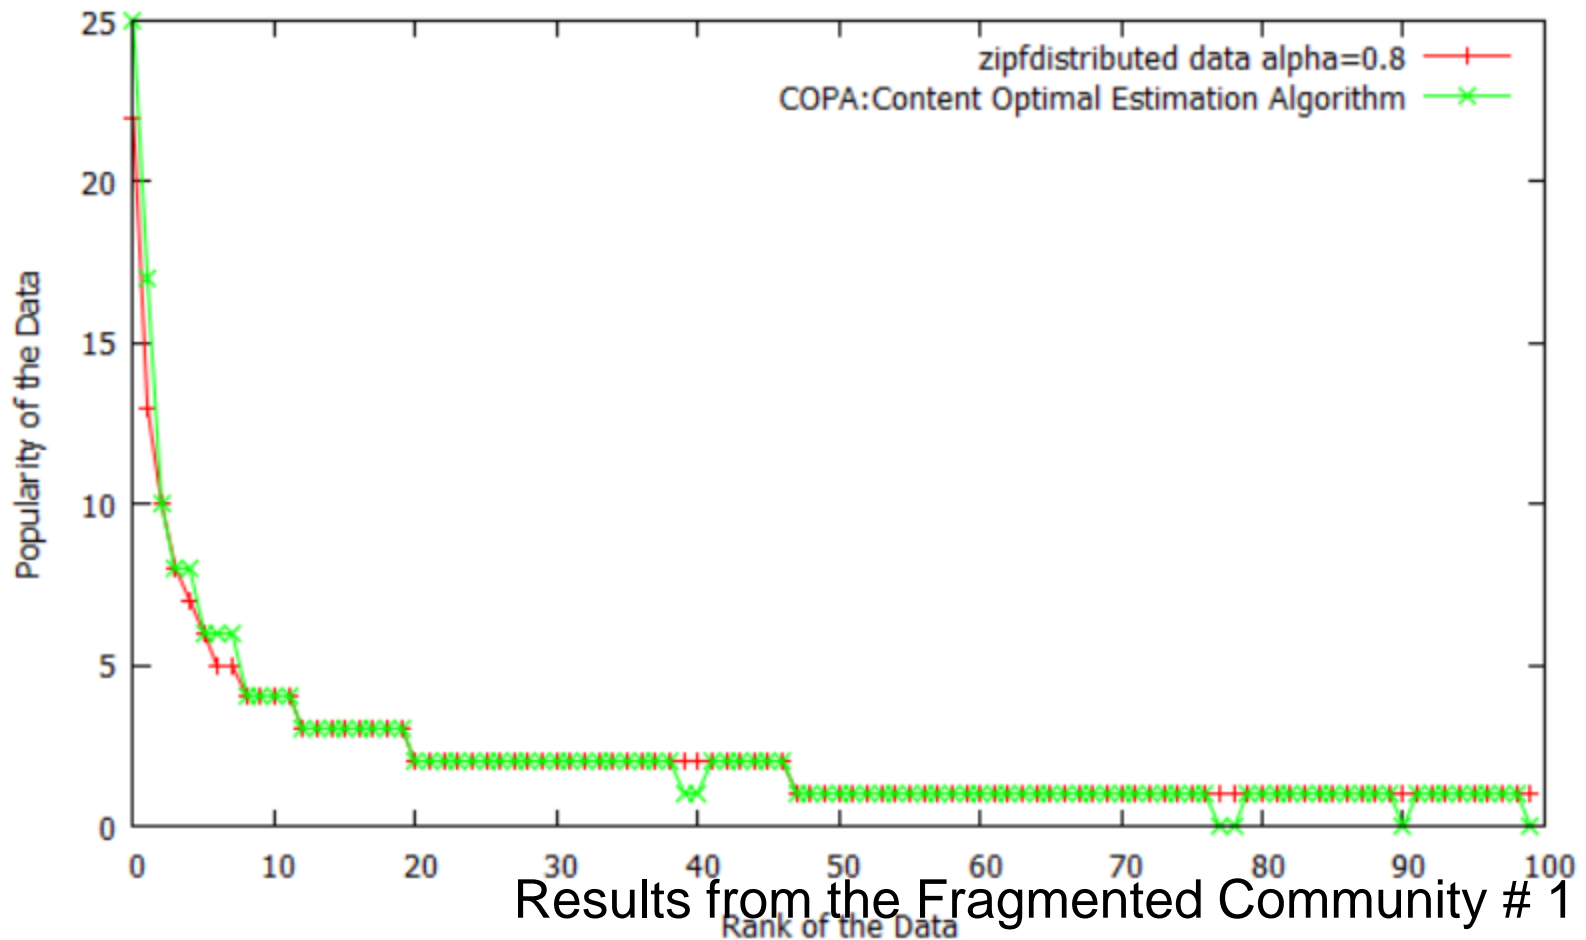
# Evaluation (15 Contacts)



Results from the Fragmented Community # 1


Zipf, alpha=0.9  
15 contacts

# Zipf-Alpha=0.8



Zipf, alpha=0.8  
15 contacts

# Better Estimation with more Contacts



Zipf,  $\alpha=0.9$   
7 and 15 contacts

# Observations

- More contacts: better accuracy
  - 15 contacts lead to quite good results
  - But: consistent underestimation (for all objects) – so may not be a real problem
- There can be more than one nonce per object name in the network
  - Can lead to slight overestimation
- PITs can become large
- Our approach significantly more memory-efficient than naive approach (thanks to aggregation)

# Summary

- **Popularity estimation in ICN** for optimizing performance and availability (here: DTN scenario)
- **Leveraging ICN naming and storage features** – changing CCN semantics (store-carry-forward of interests for long time)
- **Simple scheme – intuitively scalable**
- First evaluations suggest **sufficient accuracy**
- Next Steps
  - Polish implementation, documentation, more experiments
  - Mobile GW implementation

Empowered by Innovation

**NEC**

[www.greenicn.org](http://www.greenicn.org)