

A CONTENT-CENTRIC NETWORKING FORWARDING DESIGN FOR A NETWORK PROCESSOR

parc[®]
A Xerox Company

Marc Mosko, Palo Alto Research Center (PARC)

IEEE ICC 2015 – Next Generation Networks (NGN)

How to implement token-by-token Longest Prefix Match with variable length tokens using hardware-assisted hash tables

HARDWARE HASH TABLE

Hardware accepts keys of fixed sizes
(e.g. up to 48 bytes)

And stores data of fixed size
(e.g. up to 96 bytes)

Performance degrades with longer sizes

If a name component does not fit in the
key, we need to compress key.

RELEVANCE

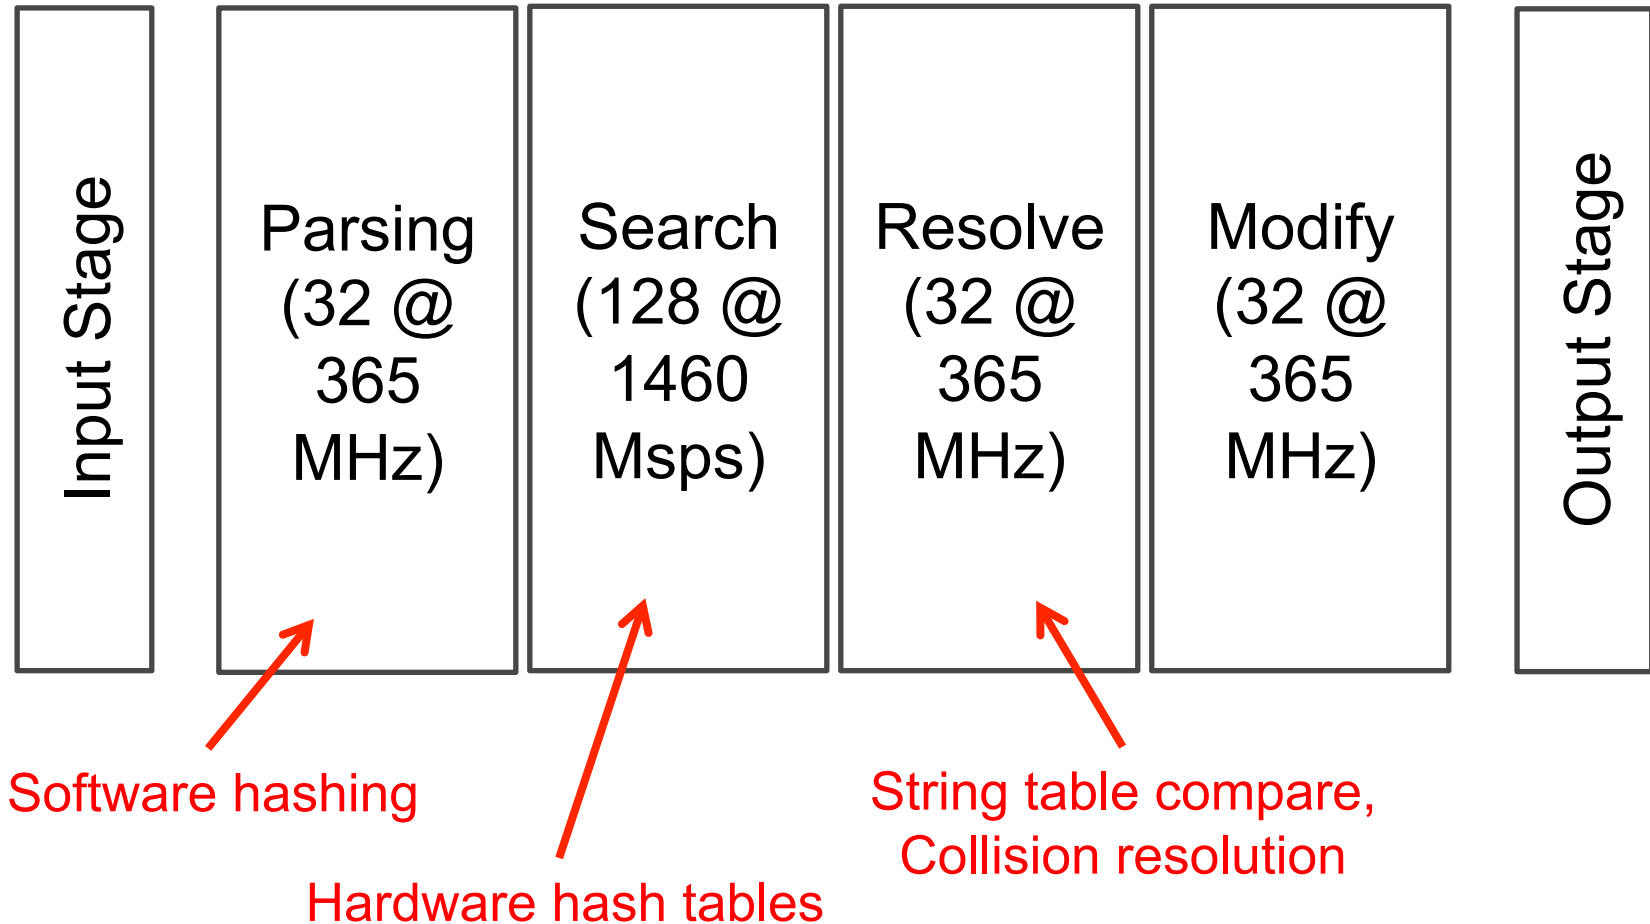
Implement Content Centric Networking
(CCNx)

On network processors
(EZchip NP4)

In multi-slot chassis switch

EZCHIP NP4

Task Optimized Processors (TOPs)



OUTLINE

ICN/CCNx Introduction
Methodology
Data structures + Algorithms
Results

INFORMATION CENTRIC NETWORKS

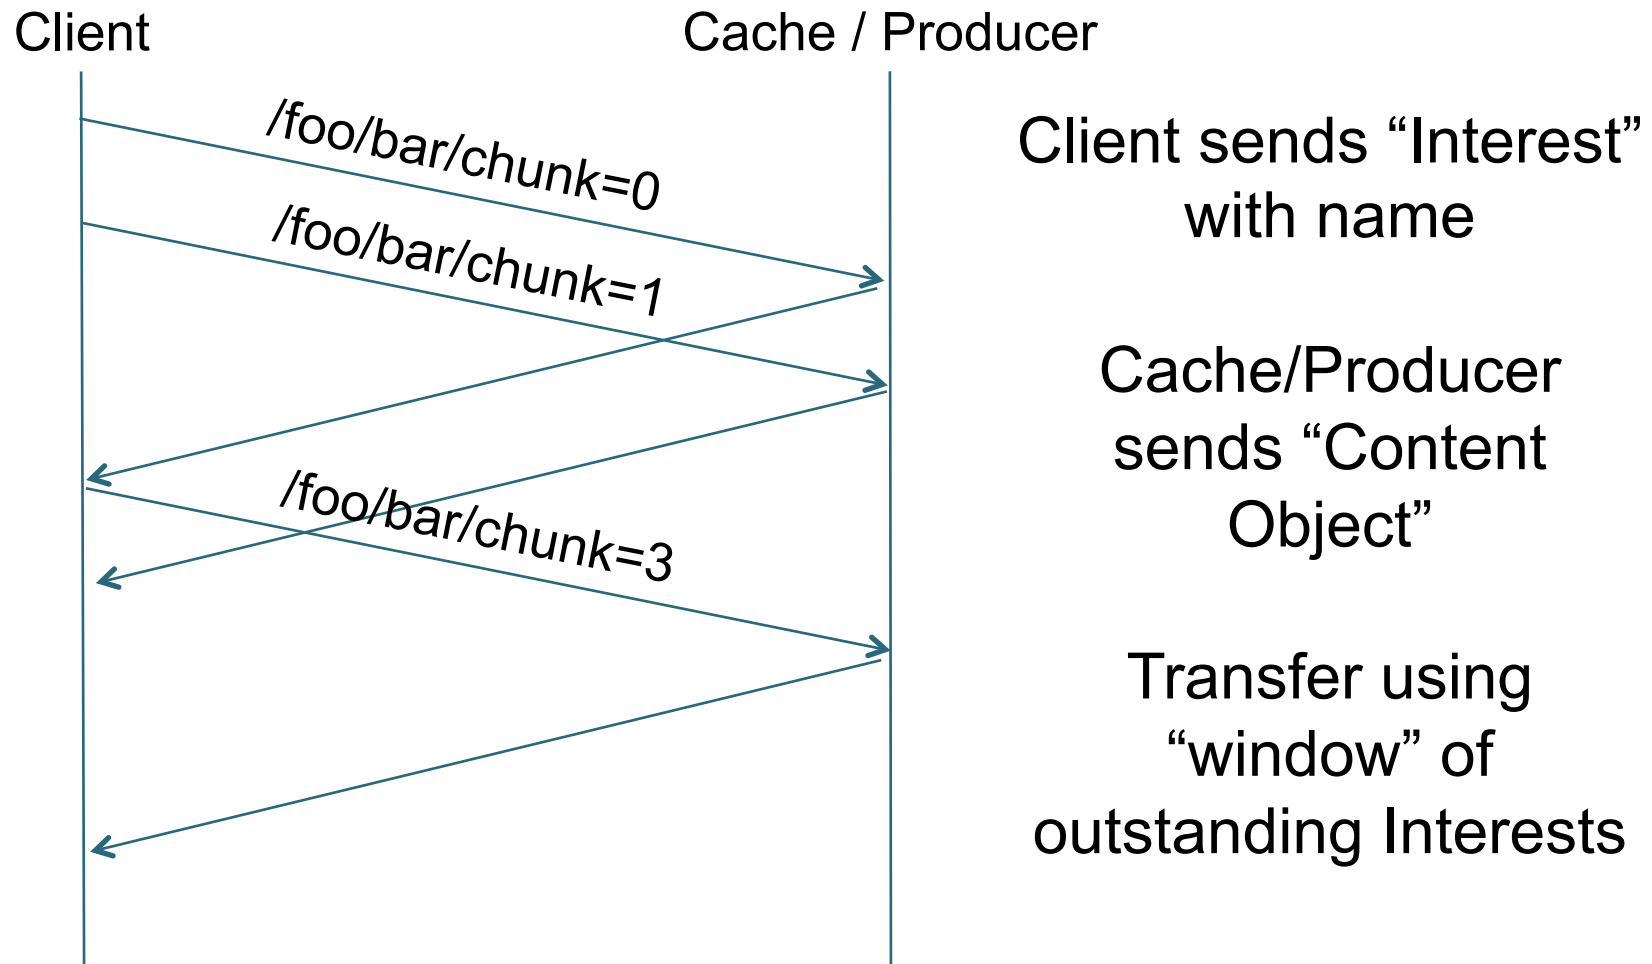
Name the data

Transfer data based on the names

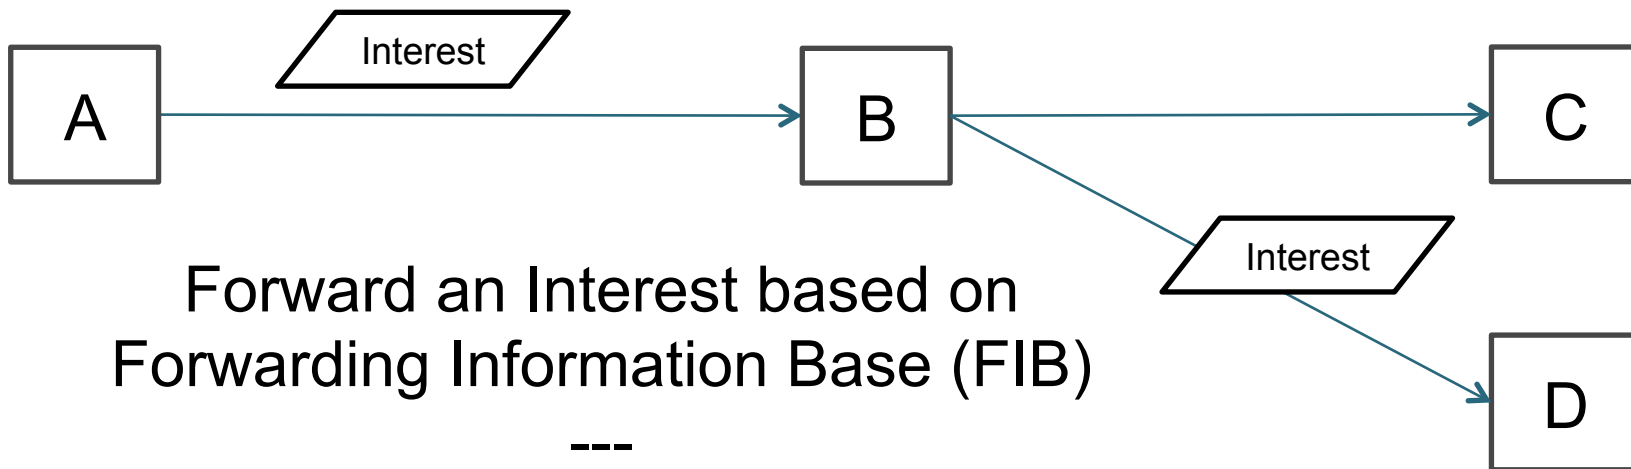
Break end-to-end paradigm

Ted Nelson's Project Xanadu (1979)

REQUEST/RESPONSE PROTOCOL



FORWARDER'S JOB



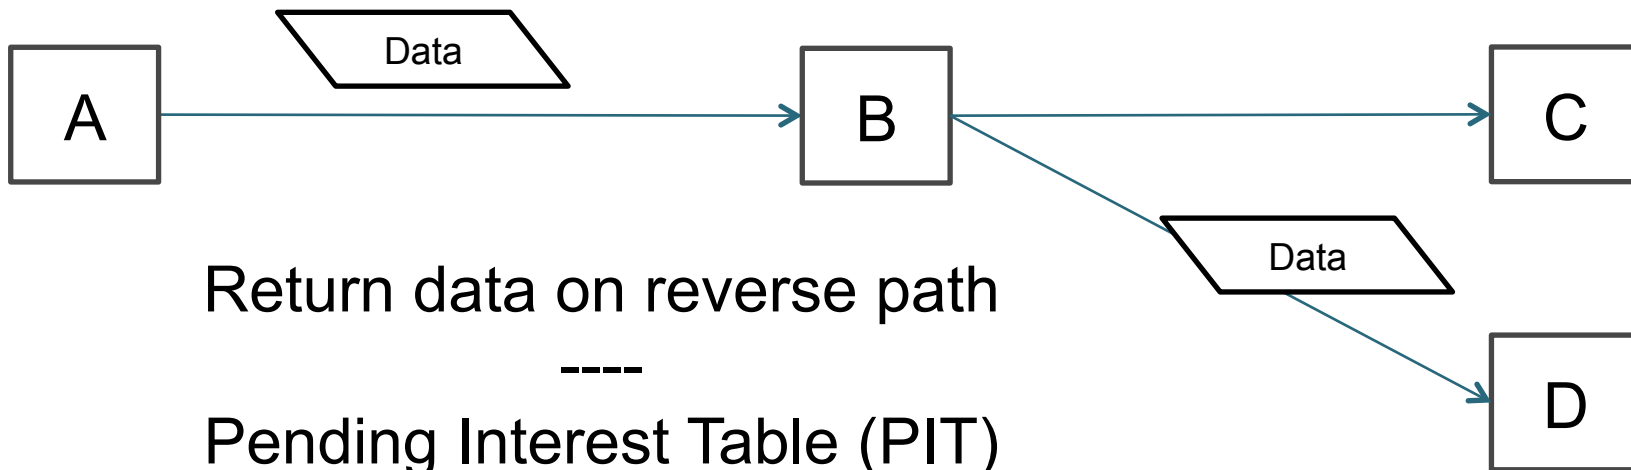
Forward an Interest based on
Forwarding Information Base (FIB)

Token-by-token Longest Prefix Match
(LPM)

lci:/apple/banana/cherry

(the topic of this talk)

FORWARDER'S JOB



(not covered in this talk)

METHODOLOGY

Statistical description of CCNx Names
+
Data Structures & Algorithms
+
Hardware Performance Model
=
Expected Performance

WHAT IS A NAME?

Use the Stanford WebBase for URIs
(<http://ilpubs.stanford.edu:8090/652>)

March 2014


64 million pages from 39,624 web sites.

Download of Links is 251GB.

Yielded 275 million unique URIs.

NAME ANATOMY

http://host.dom.tld/a/b/c/d?query_string



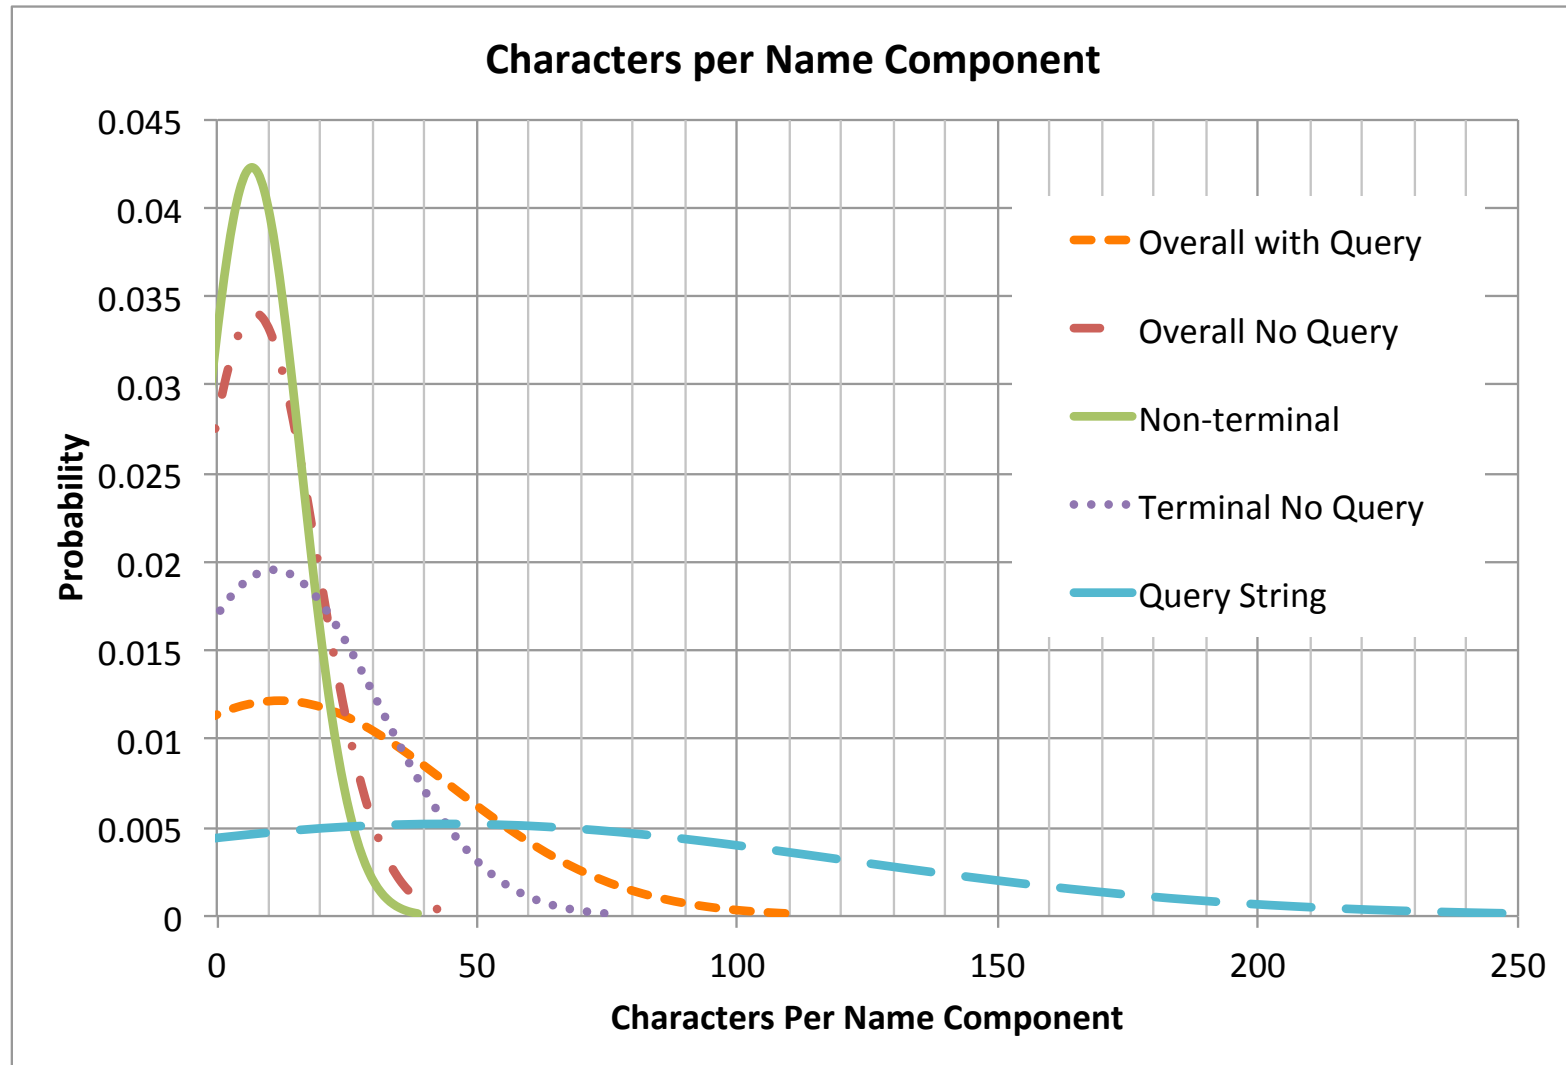
lci:/tld/dom/host/a/b/c/d/query_string



Non-terminal
tokens

Terminal
token

ANALYSIS OF URI NAMES



MAIN TAKEAWAY

	mean	stdev	99% bound
Name components	7.08	1.99	12.3

	mean	stdev	99% bound
Component Length (Overall With Query)	12.03	32.8	98.0
Component Length (Overall No Query)	7.3	11.7	38.0
Component Length (Non-terminal)	6.8	9.4	31.3
Component Length (Query String alone)	44.1	76.7	245.1

Some of these are too big for a
hardware hash table

MAIN TAKEAWAY

	mean	stdev	99% bound
Name components	7.08	1.99	12.3

	mean	stdev	99% bound
Component Length (Overall With Query)	12.03	32.8	98.0
Component Length (Overall No Query)	7.3	11.7	38.0
Component Length (Non-terminal)	6.8	9.4	31.3
Component Length (Query String alone)	44.1	76.7	245.1

The QueryString (app data) more than triples the storage requirements of the Name

ALGORITHM: INCREMENTAL HYBRID TABLES

$\text{Key} = \text{ParentKey} + \text{NameComponent}$

Or

$\text{Key} = \text{ParentKey} +$
 $\text{SWHash}(\text{NameComponent})$

EXAMPLE

T	L		T	L		T	L	
1	5	apple	1	3	pie	2	47	Abcd...

$K1 = 0x00000000 + \text{"0x00010005apple"} + \text{pad}$

➤ $EID = 0x00000011 + \text{FIB entry}$

$K2 = 0x00000011 + \text{"0x000010003pie"} + \text{pad}$

➤ $EID = 0x00002200 + \text{FIB entry}$

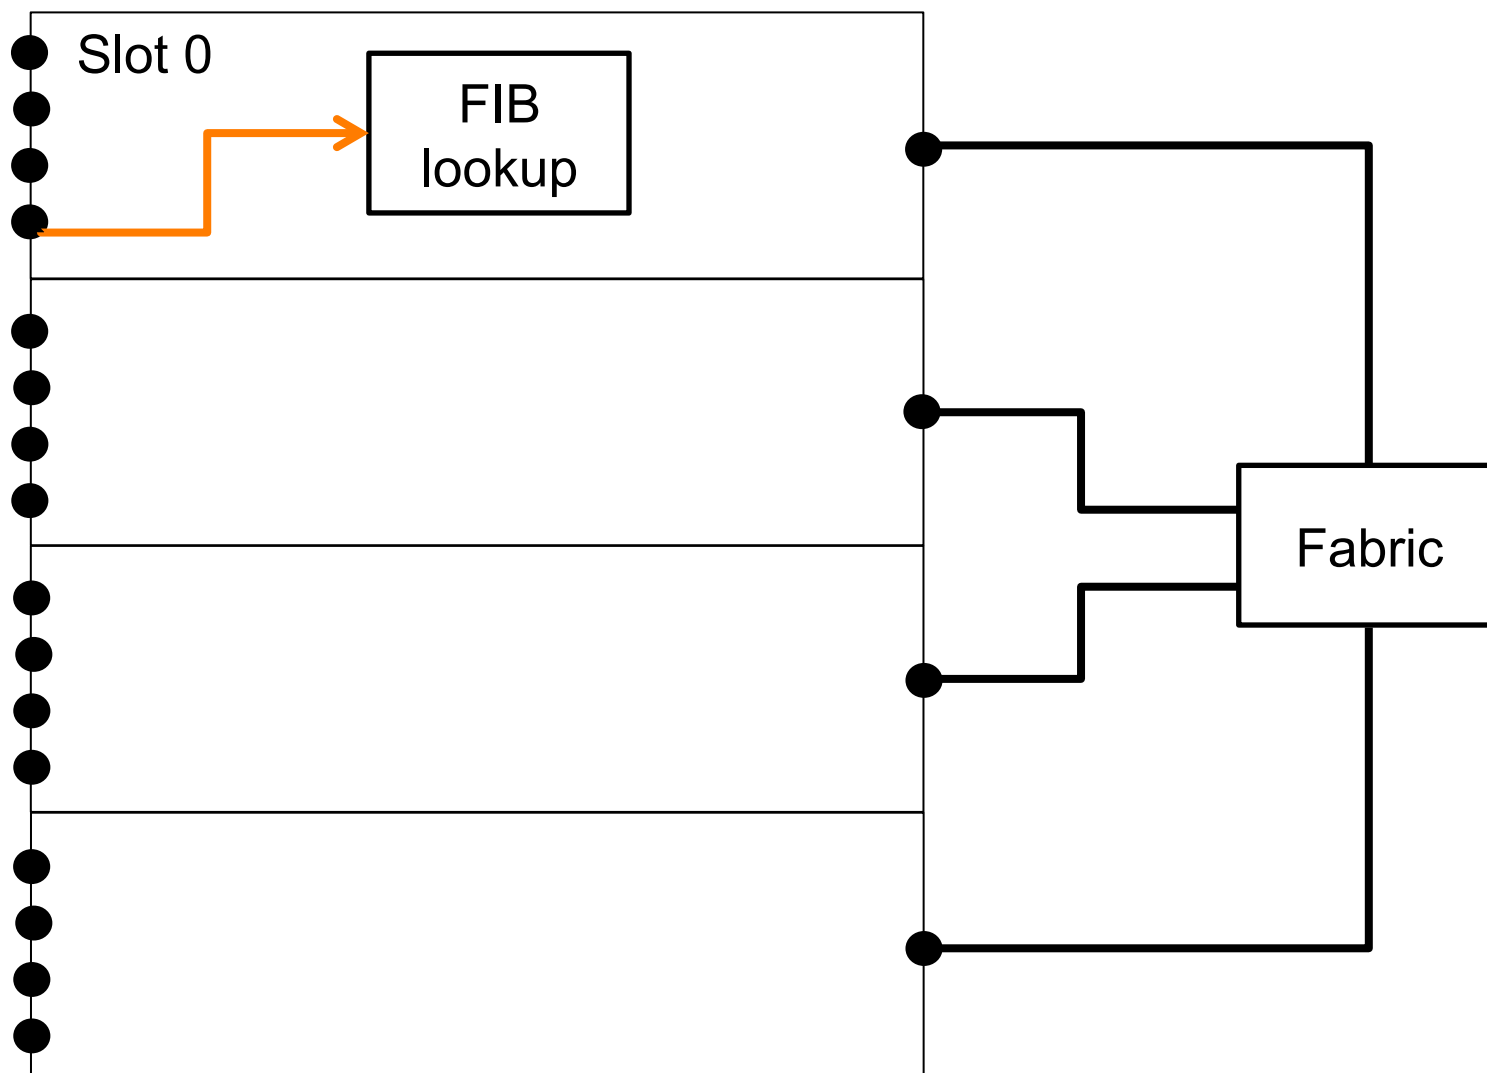
$K3 = 0x00002200 + \text{Hash("0x0002002FAbcd...")}$

➤ $EID = 0x00770000 + \text{FIB entry} + \text{String ID}$

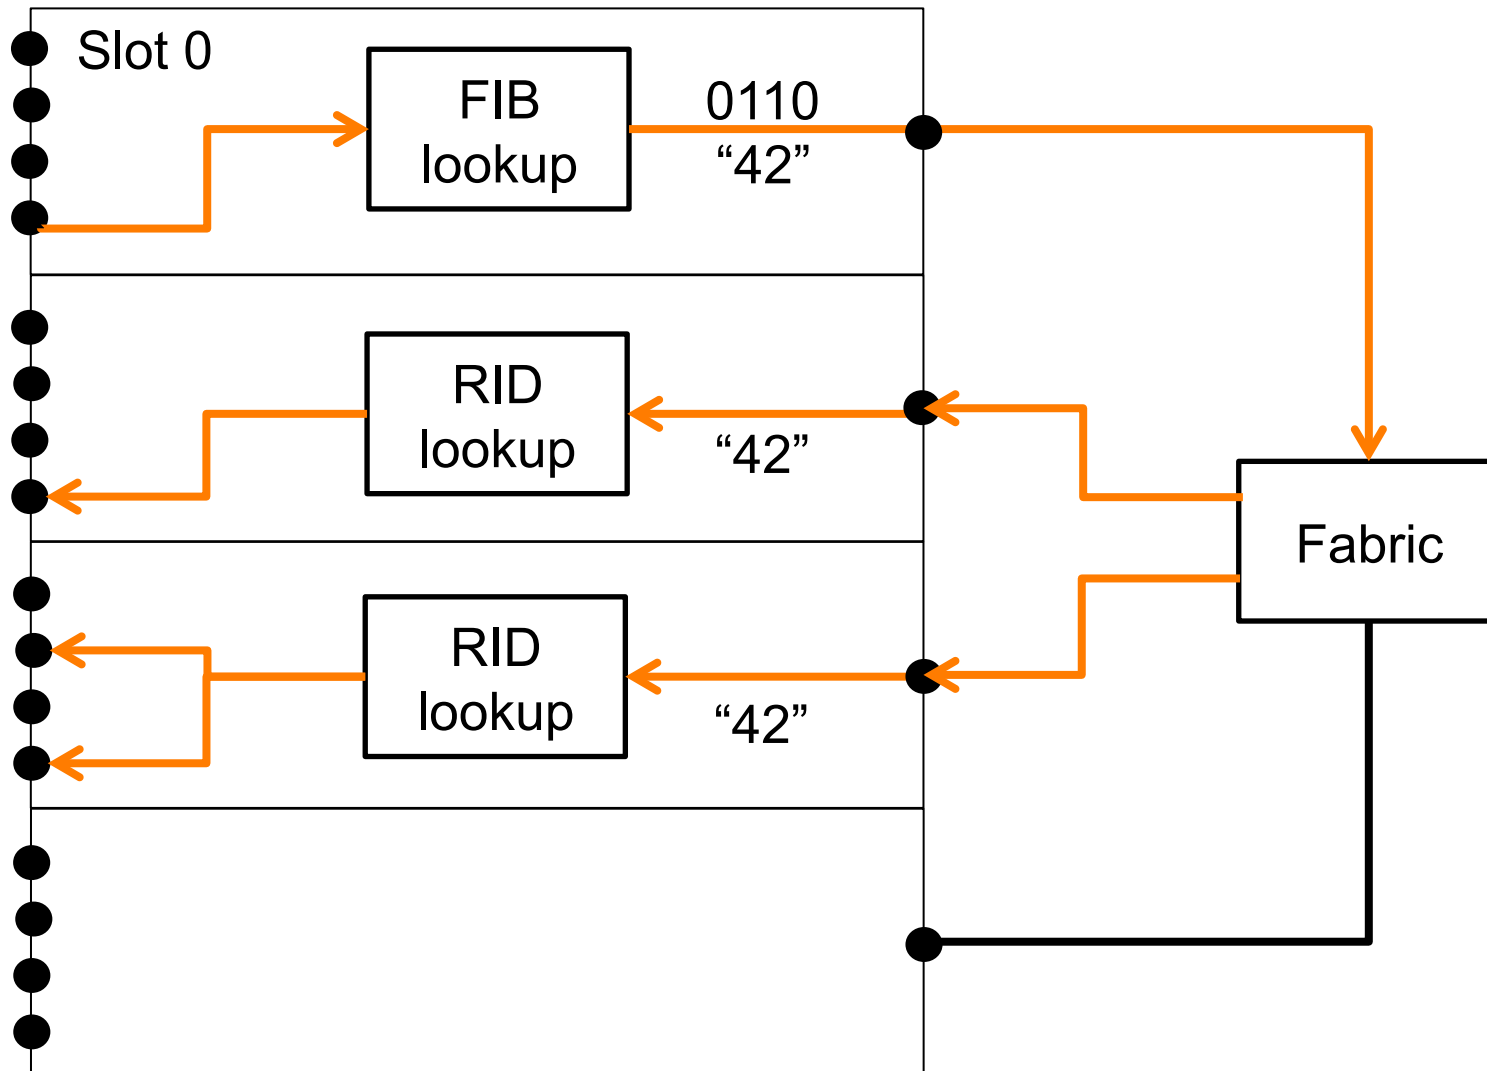
DATA STRUCTURES

- Control plane knows names $N[i]$ with components $N[i,j]$
 - It inserts FIB entries on each line card and uses a bitmap for egress card slots. It also contains a Route ID (RID) programmed on each egress slot to resolve specific media ports.
 - Fabric switching is done on the egress bitmap and carries the RID.
 - On each egress card, the RID resolves to the specific egress media ports on that card.
 - For name components that are too large for a hardware hash table key, the control plane also inserts the full name component in a string table identified by a String ID (SID).
 - If two or more name components (plus parent id) collide, there is also a Collision ID (CID), but for the hash size used this is very rare.

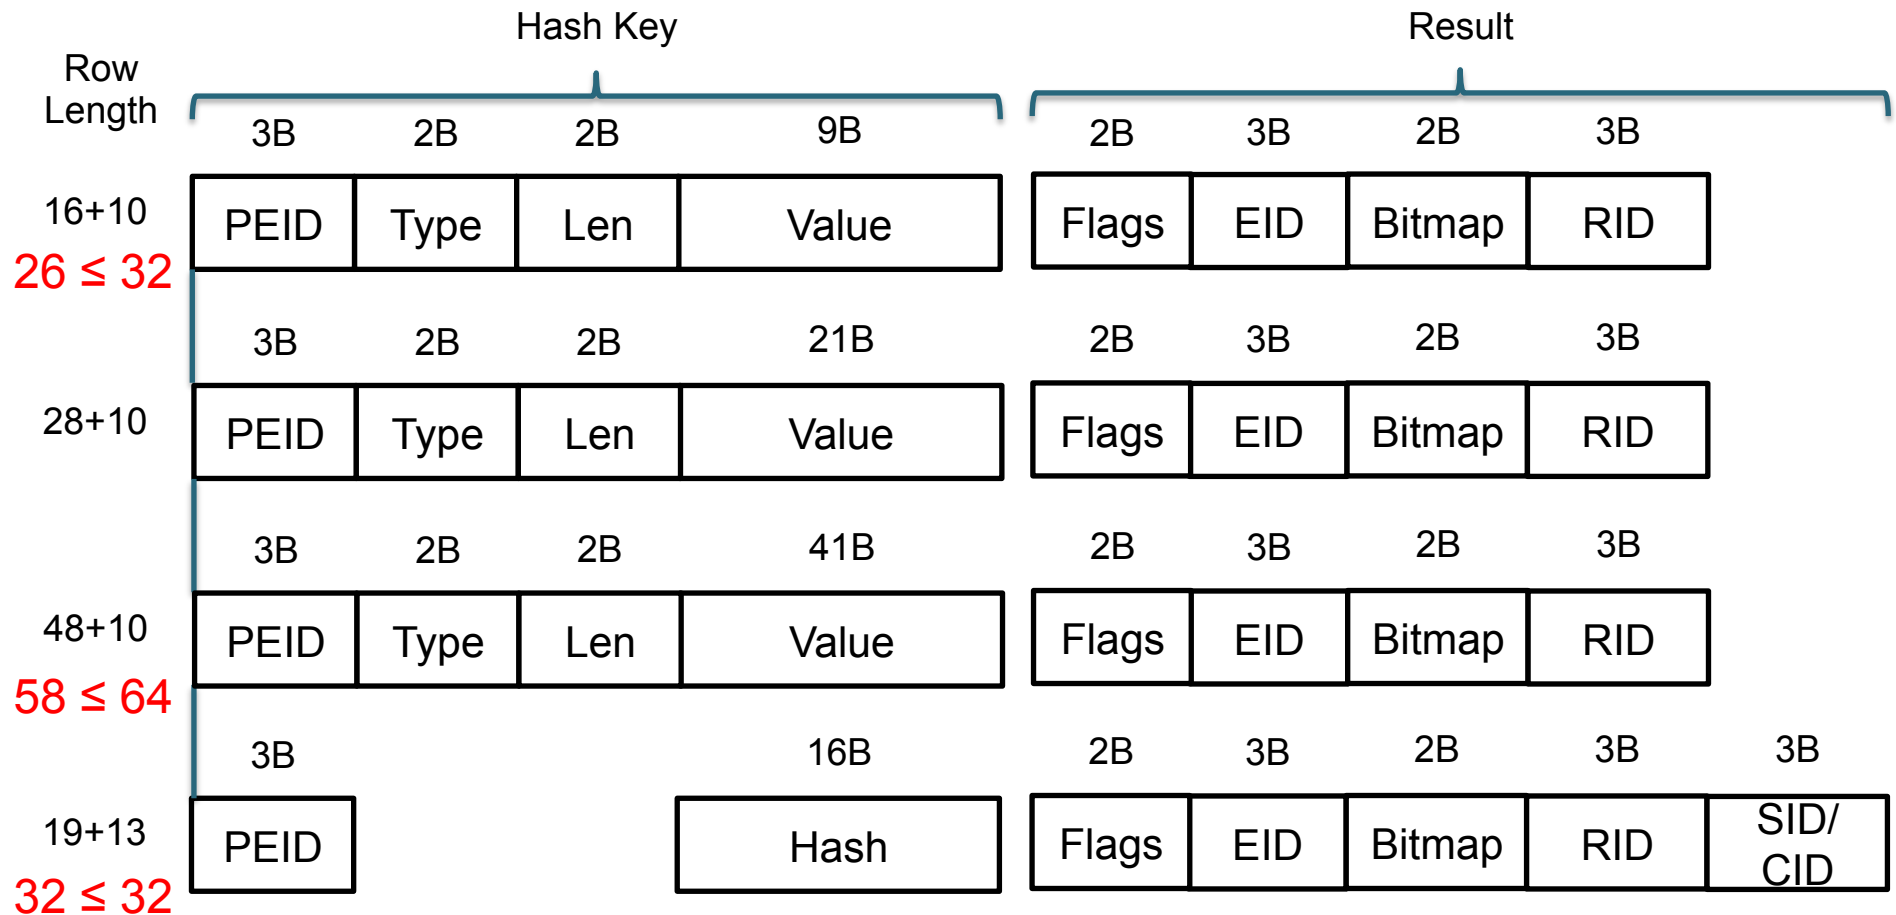
EXAMPLE



EXAMPLE



HYBRID TABLES



FIB TABLE RESULT

Flags	EID	Bitmap	RID	SID/ CID
-------	-----	--------	-----	-------------

Flags	NP4 flags
EID	Entry ID (used as Parent ID in next lookup)
Bitmap	Indicates egress slots
RID	Route ID (index to table on each egress card)
SID/CID	String ID or Collision ID (software hash lookups)

WHAT CAN GO WRONG?

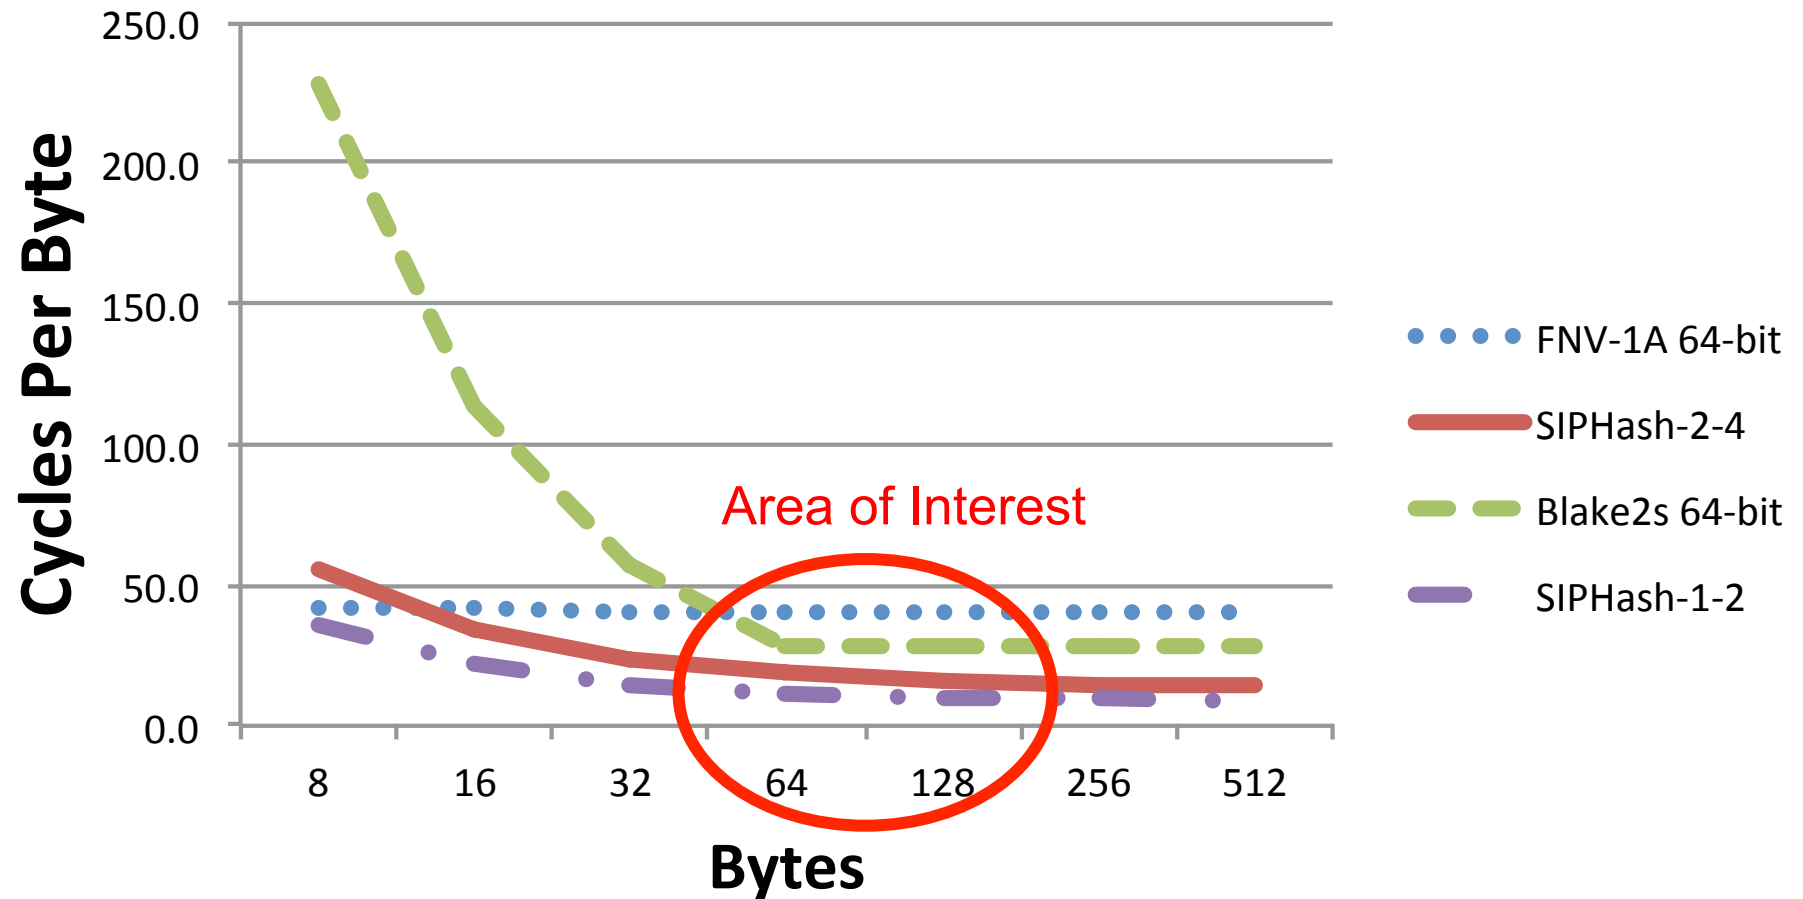
Software Hash will result in key collisions

If routing process inserts two names with same hash, it detects and indicates a “Collision ID” for second lookup

A lookup on a name may not be an actual match (hash bucket collision) so still need to do a string comparison via “String ID” lookup

SOFTWARE HASHING

Instruction Cycles per Byte



ARCHITECTURE SUMMARY

	Hardware Hash Table (up to 9 chars)

	Hardware Hash Table (up to 41 chars)

	Software Hash Table (over 41 chars)

	String Table (over 41 chars)

	RID Table

	Collision Table

TABLE I
PROBABILITY WEIGHTED BYTES AND CORE CYCLES PER ENTRY


Characters	Probability	Memory Bytes	Core Cycles
1 - 9	59.42%	21.39	76.06
10 - 41	40.57%	27.59	77.89
42+	0.01%	0.03	0.03
Average		49.00	153.97

(a) Non-terminal Components

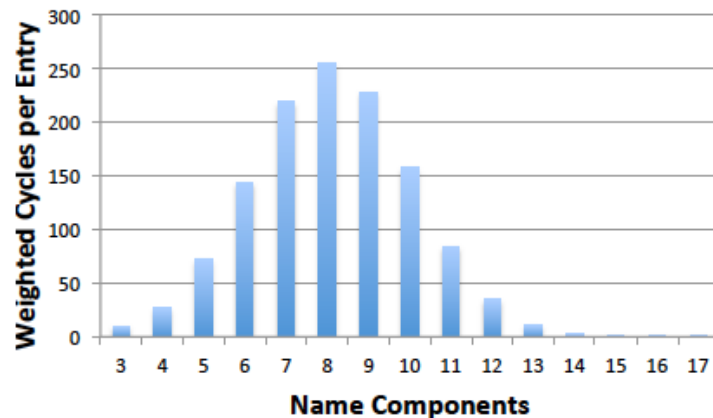
Characters	Probability	Memory Bytes	Core Cycles
1 - 9	46.34%	16.68	59.31
10 - 41	46.67%	31.74	89.61
42+	6.99%	15.94	17.90
Average		64.36	166.82

(b) Terminal Components

FINAL RESULTS



(a) Weighted Bytes



(b) Weighted Cycles

Fig. 5. Probability weighted name distributions

Summing 3 ... 17

Average FIB Entry is
378.3 bytes

Average Lookup is
1246.2 core cycles

EXPECTED PERFORMANCE

2GB DRAM stores 5.6M FIB entries

1246 core cycles over 128 lookup engines
service 37 Mpps

Computing SIPHash 2-4 for the 0.01% of
long tokens is done in the Parse TOP takes
1178 cycles, done in parallel with search

CONCLUSION

- Analysis of today's URIs shows that most names will fit in hardware hash tables.
- Implemented several hash functions on EZchip NP4.
- Propose a multi-stage incremental hash lookup to allow mixing plain keys and software compressed keys in hardware hash tables.
- Using performance model from EZchip, estimate 37Mpps.
- Software compressed keys are calculated in separate processor from hash lookups.